

Machine Learning for Biology and Health

CSCI 1851
Spring 2026

Ritambhara Singh

March 31, 2026
Tuesday



68

Today's goal – Learn about deep CNN architectures and new applications

(1) Deep CNN architectures

(2) Deep CNNs in health applications

(3) Class activity: CNNs for genomics application

(4) New dataset: Single-cell gene expression

Our problem:

Input: $\mathbb{X}$

Pixel Grid

$x^{(1)} =$



128x128 pixels

$x^{(2)} =$



→ Function: f →

$f(\mathbb{X}) \rightarrow \mathbb{Y}$

Target: $\mathbb{Y}$

Has Pneumonia?


$y^{(1)} = 0$



$y^{(2)} = 1$




CNN Architecture




Deep CNN architectures

More Complicated Networks

AlexNet:




VGG:





Can you guess what was the biggest bottleneck to adding more layers?

Revolution of Depth



Revolution of Depth


Vanishing gradients!




More Complicated Networks

ResNet:

Lots of layers, tons of learnable parameters
Avoids Vanishing Gradient problem
but how?




K. He, X. Zhang, S. Ren, and J. Sun. Deep residual learning for image recognition.
arXiv preprint arXiv:1512.03385, 2015.



Residual Blocks

- In very deep nets, each layer often needs to learn just a small transformation of the preceding layer (identity + change)
- Idea: explicitly design the network such that the output of each layer is the identity + some deviation from it
 - Deviation is known as a residual




Any questions?




Residual Blocks

- In very deep nets, each layer often needs to learn just a small transformation of the preceding layer (identity + change)
- Idea: explicitly design the network such that the output of each layer is the identity + some deviation from it
 - Deviation is known as a residual
- Allows gradient to flow through two pathways
- **Significantly stabilizes training of very deep networks**



Batch Normalization (stabilizing training)

Idea: normalize the activations for each feature at each layer




Why might we want to do this?

Batch Normalization: Motivation

More stable inputs = faster training

MNIST test accuracy vs number of training steps



<https://arxiv.org/pdf/1502.03167.pdf>

Batch Normalization: Implementation

For each feature x , Start by calculating the batch mean and standard deviation for each feature:

$$\mu_{batch} = \frac{\sum_{i=0}^{batch_size} x_i}{batch_size}$$

$$\sigma_{batch} = \sqrt{\frac{\sum_{i=0}^{batch_size} (x_i - \mu_{batch})^2}{batch_size}}$$

Batch Normalization: Implementation

Normalize by subtracting feature x 's batch mean, then divide by batch standard deviation.

$$x' = \frac{x - \mu_{batch}}{\sigma_{batch}}$$

Feature x now has mean 0 and variance 1 along the batch


Have we seen this before?

Batch Normalization in Pytorch


`nn.BatchNorm1d(num_features)`

Deep CNNs in health applications

Resnet-based model for classifying chest X-rays as normal or abnormal




Resnet-based model for classifying chest X-rays as normal or abnormal




Model	Accuracy	Recall	Precision	AUC
Custom model (ResNet50 Base)	0.7866	0.6873	0.9615	0.9023
With contrast enhancement	0.8190	0.7423	0.9600	0.9020
Custom model (ResNet101 Base)	0.7845	0.7113	0.9283	0.8823
With contrast enhancement	0.7931	0.7354	0.9185	0.8834
Custom model (ResNet152 Base)	0.7672	0.6667	0.9463	0.8664
With contrast enhancement	0.7672	0.6667	0.9463	0.8724
ResNet50 Base (no additional layers or training)	0.3685	0.0103	0.3750	0.2815
With contrast enhancement	0.3685	0.0103	0.3750	0.2812

Brain tumor segmentation




3D U-Net for brain tumor segmentation



3D Convolution


modalities

3D input



conv

3D kernel $\times 3$




=

Why 3D??


3 channels

2D output




Transposed Convolution

The transposed convolution *broadcasts* input elements via the kernel, thereby producing an output that is larger than the input



Transposed convolution with a 2×2 kernel. The shaded portions are a portion of an intermediate tensor as well as the input and kernel tensor elements used for the computation.

3D U-Net for brain tumor segmentation



Any questions?



CNNs for genomics application

How do we view the DNA

What is a gene?

Coordinates: chr1: 14569-14780

Gene

FASTA file

GGCGAGTCCGCGGGGGAGGGCGGGCAGAGCCTGAGCTCAGGTCTTTCTGCGTCTGGCGGA GGCGAGTCCGCGGGGGAGGGCGGG

```
>NG_008679.1:5001-38170 Homo sapiens paired box 6 (PAX6) 3
ACCCCTCTTTTCTTATCATTGACATTTAACTCTGGGGCAGGTCCTCGCGTAGAACGCGGCTGTCAGATCT
GCCACTTCCCCTGCCGAGCGGCGGTGAGAAAGTGTGGGAACCGGCGCTGCCAGGCTCACCTGCCTCCCCGC
CCTCCGCTCCCAGGTAACCGCCCGGGCTCCGGCCCCGGCCCGGCTCGGGGCCCCGCGGGGCCTCTCCGCTG
CCAGCGACTGCTGTCCCCAAATCAAAGCCCGCCCCAAGTGGCCCCGGGGCTTGATTTTTTGCTTTTAAAG
GAGGCATACAAAGATGGAAGCGAGTTACTGAGGGAGGGATAGGAAGGGGGGTGGAGGAGGGACTTGTCTT
TGCCGAGTGTGCTCTTCTGCAAAAGTAGCAAAATGTTCCACTCCTAAGAGTGGACTTCCAGTCCGGCCCT
GAGCTGGGAGTAGGGGGCGGGAGTCTGCTGCTGCTGTCTGCTAAAGCCACTCGCGACCGCGAAAAATGCA
GGAGGTGGGGACGCACTTTGCATCCAGACCTCCTCTGCATCGCAGTTCACGACATCCACGCTTGGGAAAG
TCCGTACCCGCGCCTGGAGCGCTTAAAGACACCCTGCCGCGGGTTCGGGCGAGGTGCAGCAGAAGTTCCC
GCGGTTGCAAAGTGCAGATGGCTGGACCGCAACAAAGTCTAGAGATGGGGTTCGTTTCTCAGAAAGACGC
```

Deep Learning for DNA sequence classification

>KM822100.1 |Lassa virus strain LASV1000-NIG-2009 segment L Z protein (Z) and polymerase (L) genes, complete cds
cccttcagtcttaggctacccccacttgact---gaatcataactctttgc-----tttgatttactatagcaaagagcataa-----
cctttgaaatcaatccattctccaagagattcaagaagttcactcagtatgtagcataatccacaggtctctgtcttgtgaggattgagacaatgtccacacctgtccattgaattctttgtt
ttaatctatcagacatcatt.....

>KM822015.1 |Lassa virus strain LASV056-NIG-2008 segment L Z protein (Z) and polymerase (L) genes, complete cds
cccttagtcttaggctgccccacttgact---gaatcataactctttgc-----ctcgatttactgtagcaaagagcatag-----
cctttgaaatcaatccattctccaagagattcaagaagctcactcagtatatcagcatgatccacaggtctctgttttgtgaggattgaaacaatgtccacacccgtccattgaatcctttgtt
ttcaatctatcggacatcattgatttcaatgcaccaactaaaagatgcggggaactatcaccaatctcccttataaacatt.....

Labels: 0 (came from location X) and 1 (came from location Y)

Help your colleague design a deep learning model for this task:

- (1) How will you represent the sequences?
- (2) What model to use?
- (3) Design a basic model with the layers covered so far.
- (4) Can we do any interpretation after the model is trained?

A = 1
T = 2
C = 3
G = 4
N = 5

Help your colleague design a deep learning model for this task:


- (1) How will you represent the sequences?
- (2) What model to use?
- (3) Design a basic model with the layers covered so far.
- (4) Can we do any interpretation after the model is trained?

1D convolution


- ① Filter weights
- ② Activations after 1 conv layer.
- ③ Pooling

	A	T	A	A	C	G	-
A	0	0	1	0	0	0	0
T	0	1	0	0	0	0	0
C	0	0	0	0	1	0	0
G	0	0	0	0	0	1	0
N	0	0	0	0	0	0	1

length of the seq.
k



seq



90%.

How do we view the DNA

What is a gene?

Can we get any other information about the gene?

Coordinates: chr1: 14569-14780


GGCGAGTCCGCGGGGGAGGGCGGGCAGAGCCTGAGCTCAGGTCTTTCTGCGTCTGGCGGA GGCGAGTCCGCGGGGGAGGGCGGG

Gene

FASTA file

```
>NG_008679.1:5001-38170 Homo sapiens paired box 6 (PAX6)
ACCCCTCTTTTCTTATCATTGACATTTAACTCTGGGGCAGGTCCTCGCGTAGAACGCGGCTGTCAGATCT
GCCACTTCCCCTGCCGAGCGGCGGTGAGAAGTGTGGGAACCGGCGCTGCCAGGCTCACCTGCCTCCCCGC
CCTCCGCTCCCAGGTAACCGCCCGGGCTCCGGCCCCGGCCCGGCTCGGGGCCCGCGGGGCCTCTCCGCTG
CCAGCGACTGCTGTCCCCAAATCAAAGCCCGCCCCAAGTGGCCCCGGGGCTTGATTTTTGCTTTTAAAG
GAGGCATACAAAGATGGAAGCGAGTTACTGAGGGAGGGATAGGAAGGGGGGTGGAGGAGGGACTTGTCTT
TGCCGAGTGTGCTCTTCTGCAAAAGTAGCAAAATGTTCCACTCCTAAGAGTGGACTTCCAGTCCGGCCCT
GAGCTGGGAGTAGGGGGCGGGAGTCTGCTGCTGCTGTCTGCTAAAGCCACTCGCGACCGCGAAAAATGCA
GGAGGTGGGGACGCACTTTGCATCCAGACCTCCTCTGCATCGCAGTTCACGACATCCACGCTTGGGAAAG
TCCGTACCCGCGCCTGGAGCGCTTAAAGACACCCTGCCGCGGGTTCGGGCGAGGTGCAGCAGAAGTTCCC
GCGGTTGCAAAGTGCAGATGGCTGGACCGCAACAAAGTCTAGAGATGGGGTTCGTTTCTCAGAAAGACGC
```


Gene expression




How would you
measure gene
expression?

RNA Sequencing


① Isolate RNA from samples




② Fragment RNA into short segments



③ Convert RNA fragments into cDNA




④ Ligate sequencing adapters and amplify




⑤ Perform NGS sequencing



⑥ Map sequencing reads to the transcriptome/genome




Predicting gene expression from DNA sequence using CNNs



Explaining model predictions

a In silico mutagenesis








b Attribution methods with backpropagation



New dataset: Single-cell gene
expression

Measuring gene expression for different conditions



~20,000 genes!
for humans

Single-cell Data



Resected
patient tumor




Dissociated
tumor cells



Single-cell RNA
sequencing



Single-cell RNA-Seq (scRNA-Seq)



Read Counts

	Cell 1	Cell 2	...
Gene 1	18	0	
Gene 2	1010	506	
Gene 3	0	49	
Gene 4	22	0	
...			

Compare gene expression profiles of single cells



Any questions?



Challenges of working with single-cell data

- Data
 - Sparsity, errors and missing data
 - “Nuisance factors”: Batch effects, etc.
- Downstream analysis:
 - Mapping to reference atlas
 - Integrating measurements


Review | [Open Access](#) | [Published: 07 February 2020](#)

Eleven grand challenges in single-cell data science

[David Lähnemann](#), [Johannes Köster](#), [...] [Alexander Schönhuth](#) 



Challenges of working with single-cell data

- Data
 - Sparsity, errors and missing data



Challenges of working with single-cell data


- Data
 - “Nuisance factors”: Batch effects, etc.




Challenges of working with single-cell data

- Downstream analysis:
 - Mapping to reference atlas

Can you train a deep learning model for this task?



Cell-type classification using deep learning



Challenges of working with single-cell data

- Data
 - Sparsity, errors and missing data
 - “Nuisance factors”: Batch effects, etc.
- Downstream analysis:
 - Mapping to reference atlas
 - Integrating measurements

Any questions?



Review | [Open Access](#) | [Published: 07 February 2020](#)

Eleven grand challenges in single-cell data science

[David Lähnemann](#), [Johannes Köster](#), [\[...\] Alexander Schönhuth](#) 

Recap – Learn about deep CNN architectures and new applications

(1) Deep CNN architectures

(2) Deep CNNs in health applications

(3) Class activity: CNNs for genomics application

(4) New dataset: Single-cell gene expression

Wrap up



What was the clearest point today?

What was the muddiest point today?

